

Deep Learning / Fall 2022

Homework 2

Please upload your assignments on or before 5pm ET on Friday October 14, 2022.

- You are encouraged to discuss ideas with each other. But you **must acknowledge** your collaborator, and you **must compose your own** writeup and code independently.
- We **require** answers to theory questions to be written in LaTeX. (Figures can be hand-drawn, but any text or equations must be typeset.) Handwritten homework submissions will not be graded.
- We **require** answers to coding questions in the form of a Jupyter notebook. It is **important** to include brief, coherent explanations of both your code and your results to show us your understanding. Use the text block feature of Jupyter notebooks to include explanations.
- Upload both your theory and coding answers in the form of a **single PDF** on Gradescope.

-
1. **(3 points)** *Designing convolution filters by hand.* Consider an input 2D image and a 3×3 filter (say w) applied to this image. The goal is to guess good filters which implement each of the following elementary image processing operations.
 - a. Write down the weights of w which acts as a *blurring* filter, i.e., the output is a blurry form in the input.
 - b. Write down the weights of w which acts as a *sharpening* filter in the horizontal direction.
 - c. Write down the weights of w which acts as a *sharpening* filter in the vertical direction.
 - d. Write down the weights of w which act as a *sharpening* filter in a *diagonal* (bottom-left to top-right) direction.
 2. **(3 points)** *Weight decay.* The use of ℓ_2 regularization for training multi-layer neural networks has a special name: *weight decay*. Assume an arbitrary dataset $\{(x_i, y_i)\}_{i=1}^n$ and a loss function $\mathcal{L}(w)$ where w is a vector that represents all the trainable weights (and biases).
 - a. Write down the ℓ_2 regularized loss, using a weighting parameter λ for the regularizer.
 - b. Derive the gradient descent update rules for this loss.
 - c. Conclude that in each update, the weights are “shrunk” or “decayed” by a multiplicative factor before applying the descent update.
 - d. What does increasing λ achieve algorithmically, and how should the learning rate be chosen to make the updates stable?
 3. **(2 points)** *The IoU metric.* In class we defined the IoU metric (or the Jaccard similarity index) for comparing bounding boxes.
 - a. Using elementary properties of sets, argue that the IoU metric between any two pair of bounding boxes is always a non-negative real number in $[0, 1]$.
 - b. If we represent each bounding box as a function of the top-left and bottom-right coordinates (assume all coordinates are real numbers) then argue that the IoU metric is *non-differentiable* and hence cannot be directly optimized by gradient descent.
 4. **(4 points)** *Training AlexNet.* Open the (incomplete) Jupyter [notebook](#) in Google Colab (or other Python IDE of your choice) and complete the missing items.

5. **(3 points)** *Object detection.* We will be finetuning a pretrained torchvision object detection model in this problem. It will be helpful to go through the excellent tutorial (with example code) [here](#). You can reuse whatever code you like from this tutorial with *proper attribution*!
- a. The tutorial shows how to finetune a pretrained model (what they call option 1). Now, follow their approach to add a different backbone (option 2). You can use the same dataset (and code) as the tutorial.
 - b. How do the performance of the two models compare on the training data after 10 epochs?
 - c. Test both models on [this](#) image of the Beatles. Be careful to make sure the dimensions align and the scaling is similar. How do the two models compare in terms of performance on this image?